


(a) Backend HPA response under controlled CPU load



Source: `kubectl get hpa` sampled every 15s; load generated by an in-cluster curl Pod.